

COMPX527 Assignment 1 Report

Glenn Cumming

Department of Computer Science
University of Waikato
Hamilton, New Zealand
glenn@hif.nz

Mitchell Grout

Department of Computer Science
University of Waikato
Hamilton, New Zealand
mjg44@students.waikato.ac.nz

Shufen Li

Department of Computer Science
University of Waikato
Hamilton, New Zealand
sl302@students.waikato.ac.nz

Sunny Chikara

Department of Computer Science
University of Waikato
Hamilton, New Zealand
sc420@students.waikato.ac.nz

YingJun Huang

Department of Computer Science
University of Waikato
Hamilton, New Zealand
yh320@students.waikato.ac.nz

Abstract—We designed an Object Detection web service on the AWS Cloud Provider, using Terraform for programmable infrastructure, Ansible for deployment, and Python for the web service and object detection.

Index Terms—AWS, COMPX527, Terraform, Ansible, Flask, Python, Object Detection, Retinanet

I. SOLUTION SUMMARY

The Flask Object Detection service is a load balanced HTTP service deployed in AWS EC2, using the AWS ELB load balancer. The service can accept a JPG image and run object detection on it, and returns the resulting image with identified objects highlighted.

II. MOTIVATION

We decided on object detection in the cloud in order to learn and appreciate the challenges of providing a stateless, scalable service that could be used for a variety of other projects. Examples included

- Contextual threat analysis, such as humans detected in an area which should only include livestock.
- Numerical analysis, such as the number of a species of wildlife in an area over time.
- Absence detection, identifying objects that should be in an image but are not

III. PROPOSED SOLUTION

An Amazon Web Services based cluster of web servers taking images from clients and returning images with bounding boxed identified objects plus json formatted data, behind an Elastic Load Balancer. The cluster will be horizontally scalable by adding web servers, and accessible via curl or any other HTTP client that can perform HTTP POST requests. It was trained using the AWS COCO data set ¹, and also used the set to for testing and demonstrations.

¹<https://registry.opendata.aws/fast-ai-coco/>

IV. SOLUTION ARCHITECTURE

We decided to use a load balanced web service architecture, a very common and well understood model with common solutions to scalability and uptime. We have the following components and features:

A. The Image Processing Server

The project uses the Retinanet based Object Detection developed by Keras ². The web service was developed in Python using a Flask server ³.

The server receives via HTTP JPG formatted image, and returns the image with bounding boxes around identified objects, plus JSON data of the images.

B. The Load Balancer

The object detection being inherently slow forces the use of load balancing and scaling techniques. Using the AWS Elastic Load Balancer allowed us to balance the HTTP requests over two or more EC2 servers running the Flask servers. We could create as many Flask servers as we wished, although of course this lent itself to wasted compute resources. Though we had great success early on using AWS AutoScaling Groups, which allows a lower and upper limit of EC2 instances running our Flask server to be defined, this functionality was not available in the student accounts.

C. The Web User Interface

In order to have a better demo, it was decided to go ahead and produce a web site interface that would allow uploading

D. The Web API

The intended standard way to interact with the Flask Object Detection Cluster is to access it via HTTP calls using the url `http://<load_balancer_fqdn>/detect`. As the system is intended for the use of non-private images no ssl is implemented on the load balancer, though it is supported.

²<https://github.com/fizyr/keras-retinanet>

³<https://pypi.org/project/Flask/>

E. Amazon Web Services Cloud

The service running on the Amazon Web Services Cloud, the required choice of the assignment. Though later in lectures there it was mentioned we could look in to using other Cloud Service Providers, we decided we had already made enough progress to commit ourselves to AWS for this project. Though we did use terraform, which supports IaaS deployments with multiple CSP, in case we had a need to effect a change quickly.

V. TECHNOLOGY

The following are the significant technologies we used in this assignment.

A. Retinanet

Retinanet is an open source object detection neural net. This technology was chosen for the following reasons ⁴.

- Free Open Source project obviously beneficial for a project with a limited budget.
- Python 3 for rapid development and deployment. With a mix of experienced and less experienced team members this was particularly useful.
- Specific instructions using it with the dataset available on the Registry of Open Data on AWS ⁵

B. Terraform

Terraform was given in as an example of the provisioning service to use in this assignment. AWS CloudFormation was another option suggested. The decision to go with terraform rather than CloudFormation was based on the following

- CloudFormation is proprietary and is specific to AWS cloud offerings. Though this assignment is on a small and temporary scale, we still not want to use a technology that would create vendor lockin
- Terraform supports provisioning many different platforms, both open standards and proprietary, such as Azure, Google Cloud, Kubernetes and OpenStack ⁶. Developing experience in deployments with terraform therefore was deemed to be more useful.

C. Ansible

Though Terraform is capable of running commands post-install in order to install services and other software needed for our cluster, our cluster used Ansible playbooks instead. Though it meant learning another technology, we deemed a good use of our time since:

- Terraform can only run scripts, which would have to be created.
- Ansible's language is very flexible and is created specifically for the purpose of deployment.
- This is a recommended approach by HashiCorp, the developers of Terraform. ⁷

⁴

⁵<https://registry.opendata.aws/fast-ai-coco/>

⁶<https://aws.amazon.com/cloudformation/>

⁷<https://www.hashicorp.com/resources/ansible-terraform-better-together>

D. Github

We used github for source control of all documents and code. This was chosen over other solutions such as SVN and Mercurial due to the familiarity of some of the team members with git. Others in our team have learnt cloning, pulling and branching. ⁸

E. Slack

As we needed to communicate effectively over a period of weeks without seeing each other often, we decided to use the collaboration software Slack ⁹. We used a free account, and created a channel to allow us to privately communicate about our work. It was used extensively and allowed constant effective communication.

F. Trello

In order to break up and manage tasks we created a Kanban Board in Trello ¹⁰. This got us started with basic learning and setup tasks early on, but as the project progressed we tended to move to the use of Slack to define and assign tasks and workloads.

G. LaTeX

LaTeX ¹¹ is the accepted standard for scientific papers; therefore although the assignment pointed us to the IEEE A4 standard for reporting, we opted to use the LaTeX standard. This gave us very useful experience in using LaTeX for proper report formatting. This report was created in LaTeX using the IEEEtran document class provided by IEEE template available at ¹², and generated into PDF by Gnome LaTeX ¹³

H. Flask

The decision to develop and deploy a service using the keras-retinanet Python library immediately suggested to us we should again use Python to provide the web service. Our team had developers with experience in both Tornado and Flask; Flask was used as the web service development fell on the person with Flask experience. Both solutions would have been suitable.

VI. EXISTING OBJECT DETECTION SOLUTIONS

There is a wide range of cloud and non cloud solutions. A comparison of these to our own solution is beyond the scope of this document. Those provided by the large Cloud Service Providers includes:

A. Rekognition

Amazon's object detection solution ¹⁴.

⁸<https://github.com/>

⁹<https://slack.com/>

¹⁰<https://trello.com/>

¹¹<https://www.latex-project.org/>

¹²<https://www.ieee.org/conferences/publishing/templates.html>

¹³<https://wiki.gnome.org/Apps/GNOME-LaTeX>

¹⁴<https://aws.amazon.com/rekognition/>

B. Vision AI

Google's object detection solution ¹⁵.

C. Computer Vision

Microsofts object detection solution ¹⁶.

VII. SECURITY AND VUNLERABILITY ASSESSMENT

A. Data Security

The Flask Object Detection Cluster was specifically created to be stateless and public. Interception of HTTP requests and results can be easily intercepted by third parties with access to the networks between the client and the service. Users of the service should be aware that any images sent or data received from the service could be illicitly obtained, maliciously modified, or corrupted in transit.

Mitigation could be achieved by adding SSL encryption and offloading it at the load balancer.

B. Access Security

a) *Identity and Access Management*: Management access of the resources residing on AWS is protected via the AWS Identity and Access Management (IAM) system ¹⁷. We used individual accounts provided by the course. If an account was compromised the it would allow malicious actors to shut down services, or replace the current services with malicious alternatives. For example, they could replace an instance with one that will return malware embedded in an image.

Mitigation can be using good protection practices for authentication information, Two Factor Auth (2FA), regular monitoring of the service with a set of well known HTTP requests checked against expected results.

b) *Server Access*: The Flask servers are publically accessible via SSH, needed for deployment. If the key pair is compromised or the keys obtained then root access to the Flask servers could be obtained. This would grant the attacker full access and the ability to use the server's resources for their own purposes.

Migration can be using ssh key passwords, encrypted storage of the keys, checking server hashes, and keeping ssh software up to date with security patches.

C. Network Security

We used AWS Security Groups to restrict the Load Balancer to only ingress traffic port 80. The Image Processors restricted to port 22 and 80.

Allowing direct access to the service port is a known problem that could allow an attacker to easily DDOS a single server, bypassing the Load Balancer. To mitigate this we need to work out how to define the CIDR range in the security group to match the internal IP of the Load Balancer. The service is load balanced, but we were not able to use AutoScaling Groups due to our account restrictions. This results in having

to manually create additional servers to cope with increased requests, and so a DDOS would be more effective since a human will have to respond to any such attack.

To mitigate this we should move to an account that allows the use of Auto Scaling Groups.

D. Monitoring and Security

AWS EC2 instance monitoring was enabled during the terraform creation of the cluster. ELB monitors the availability of the Flask HTTP servers. The lack of external monitoring means that we are totally dependent on AWS for security checks and uptime monitoring.

Mitigation of the lack of monitoring is obviously achieved by using both in house and monitoring services. Common examples of in house monitoring is via Nagios, Icinga2, Munin and Prometheus. Monitoring services include Uptime.com ¹⁸ for ping and HTTP monitoring, and Paessler ¹⁹, which provides additional services such as port monitoring.

VIII. ACTUAL AWS EXPENDITURE

TODO

IX. FUTURE IMPROVEMENTS

a) *Packaging*: Deployment scripts were chosen for deployment and running the services. However, if the service was to be developed further we would be using stdeb ²⁰ for packaging for the Ubuntu servers, and potentially set up a Personal Package Archive ²¹. This would ensure further development and upgrades by making it part of the standard apt package management system.

b) *Service Management*: Currently the starting and stopping of the service is done via control scripts. These would be integrated into systemd service management ²².

c) *HTTP*: The service is stateless, public and insecure. If a layer of security is desired to protect the data in transit, then HTTPS can be added to the Elastic Load Balancer via the AWS Certificate Manager ²³. The simple reason for not using this was cost: we would have to pay for extra services when our budget was tight.

d) *Cloudflare Multi Region Load Balancing*: Redeployment of the service changes the FQDN in the URL. Also, the service is currently deployed on only one region. The Cloudflare Load Balancing would allow the use of many regions or in fact many Cloud Providers to be used with a single content url presented to the end consumer. The fact that the service is stateless makes it very suitable to this form of scale out ²⁴

¹⁸<https://uptime.com/uptime-monitoring>

¹⁹https://www.paessler.com/port_monitoring

²⁰<https://pypi.org/project/stdeb/>

²¹<https://help.ubuntu.com/community/PPA>

²²<https://wiki.debian.org/systemd>

²³<https://docs.aws.amazon.com/elasticloadbalancing/latest/classic/ssl-server-cert.html>

²⁴<https://www.cloudflare.com/load-balancing/>

¹⁵<https://cloud.google.com/vision/>

¹⁶<https://azure.microsoft.com/en-us/services/cognitive-services/computer-vision/>

¹⁷<https://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>

e) *Monitoring*: Future development of the cluster would likely use additional monitoring of performance, costs and uptimes using dedicated EC2 instances, and such existing monitoring systems as Icinga2, Prometheus and Munin. Uptime monitoring would of necessity be run externally, such as on internal machines or another cloud providers offering.

We would also start to use third party external monitoring services such as Uptime.com ²⁵ or Paessler ²⁶

X. TEAM MEMBERS CONTRIBUTIONS

XI. ASSIGNMENT REQUIREMENTS COMPLETION

The requirements for this Assignment were communicated in several ways, including the assignment documentation, messages in the Moodle channel, the lectures, and by asking the lecturer and assistant directly. We treated these as customer requirements, which typically start with a good understanding of the desired outcomes and often require feedback from the customer to ensure we are achieving them.

The following is what we have determined to be the requirements, and how we have achieved each one.

- Use of AWS to create a non-trivial service, demonstrating our understanding of cloud technologies, their use, and their challenges.
- Use of technologies such as Terraform, Ansible and Cloud Formation for programmably defined infrastructure and deployments
- Use of an open dataset from AWS
- Work within a team with very diverse technical backgrounds and manage them to ensure fair distribution of the work.

²⁵<https://uptime.com/uptime-monitoring>

²⁶<https://www.paessler.com/>